

```
In [1]: import tensorflow as tf
from tensorflow import keras
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import pydot, graphviz
import numpy as np
import pandas as pd
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras import models, layers, optimizers
from tensorflow.keras.applications import VGG16
```

```
In [2]: # save the relevant directories containing the pictures
# put an 'r' because we want forward slash instead of backward slash
test_dir = r'C:\Users\tanch\OneDrive\Documents\Udemy\CNN for computer vision\Learning p
train_dir = r'C:\Users\tanch\OneDrive\Documents\Udemy\CNN for computer vision\Learning
validation_dir = r'C:\Users\tanch\OneDrive\Documents\Udemy\CNN for computer vision\Lear
```

## Data preprocessing

```
In [3]: # create ImageDataGenerator object
# specify to rescale pixel values to [0,1]
train_datagen = ImageDataGenerator(rescale = 1./255)
test_datagen = ImageDataGenerator(rescale = 1./255)

# this generator generates data      # directory variable specified earlier
train_generator = train_datagen.flow_from_directory(train_dir,
                                                    # photos are resized to 150 x 150
                                                    target_size = (150,150),
                                                    # IMPT: images are generated in bat
                                                    batch_size = 20,
                                                    # specify binary classification
                                                    class_mode = 'binary')

# directory variable specified earl
validation_generator = test_datagen.flow_from_directory(validation_dir,
                                                         # photos are resized to 150 x 150
                                                         target_size = (150,150),
                                                         # photos will come in batches of 20
                                                         batch_size = 20,
                                                         # specify binary classification
                                                         class_mode = 'binary')
```

Found 2000 images belonging to 2 classes.  
Found 1000 images belonging to 2 classes.

## Building CNN

```
In [4]: # specifying CNN model architecture
model = models.Sequential()
        # using 32 filters of size 3x3 NOTE: default stride=1, default
model.add(layers.Conv2D( 32, (3,3), activation='relu',
        # images were resized to 150x150 earlier
        # 3 because images are colored
        input_shape=(150,150,3)))
# each convolution layer has its own max pooling layer
model.add(layers.MaxPooling2D((2,2)))

        # 64 filters of size 3x3
model.add(layers.Conv2D( 64, (3,3), activation='relu',input_shape=(150,150,3)))
model.add(layers.MaxPooling2D((2,2)))

        # 128 filters of size 3x3
model.add(layers.Conv2D( 128, (3,3), activation='relu',input_shape=(150,150,3)))
model.add(layers.MaxPooling2D((2,2)))

        # 128 filters of size 3x3
model.add(layers.Conv2D( 128, (3,3), activation='relu',input_shape=(150,150,3)))
model.add(layers.MaxPooling2D((2,2)))

# flatten to single array
model.add(layers.Flatten())

# 512 neuron hidden layer
model.add(layers.Dense(512, activation='relu'))

# 1 neuron output layer for binary classification
# sigmoid activation for binary classification
model.add(layers.Dense(1, activation='sigmoid'))
```

```
In [5]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 512)	3211776
dense_1 (Dense)	(None, 1)	513
=====		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		
=====		

```
In [6]: model.compile(loss='binary_crossentropy',  
                      # RMSprop used over sgd # Learning rate set to 0.0001  
                      optimizer= optimizers.RMSprop(lr=1e-4),  
                      metrics=['acc'])
```

```
In [7]: # training data comes from the train_generator
history = model.fit_generator(train_generator,
                             # IMPT: 20 x 100 = 2000 images in training set
                             steps_per_epoch=100,
                             epochs=20,
                             # validation data comes from the validation generator
                             validation_data=validation_generator,
                             # IMPT 20 x 50 = 1000 images in the validation set
                             validation_steps = 50)
```

WARNING:tensorflow:From <ipython-input-7-1dd1c2e963cd>:9: Model.fit\_generator (from tensorflow.python.keras.engine.training) is deprecated and will be removed in a future version.

Instructions for updating:

Please use Model.fit, which supports generators.

WARNING:tensorflow:sample\_weight modes were coerced from

```
...
to
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...
to
['...']
```

Train for 100 steps, validate for 50 steps

Epoch 1/20

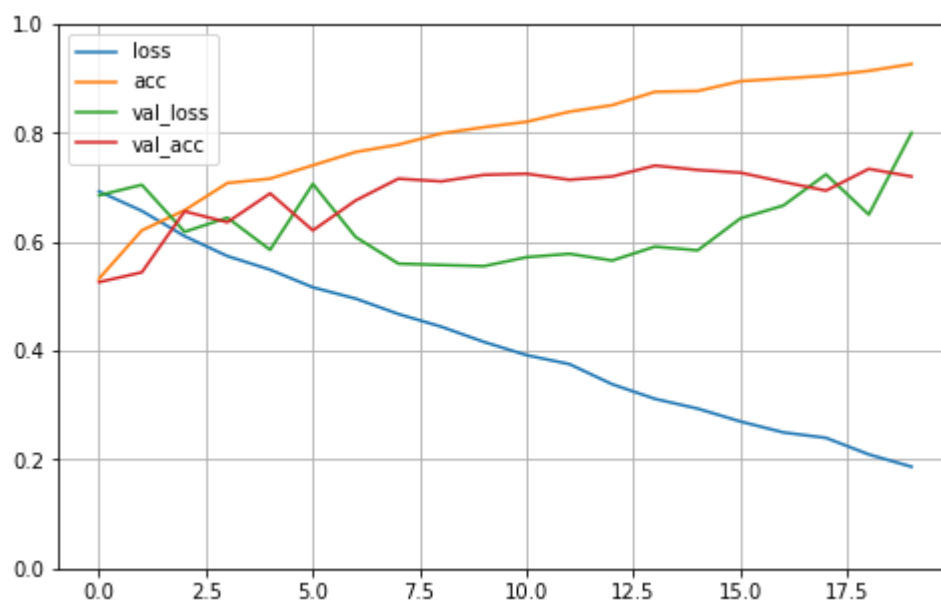
100/100 [=====] - 67s 667ms/step - loss: 0.6925 - acc: 0.5

325 - val\_loss: 0.6854 - val\_acc: 0.5260

Epoch 2/20

100/100 [=====] - 65s 648ms/step - loss: 0.6568 - acc: 0.6

```
In [8]: # visualise the training progression
pd.DataFrame(history.history).plot(figsize=(8,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.show()
```



- evident **overfitting** as seen from the increasing difference between training and validation **accruacy**

```
In [9]: # save model
model.save('CNN_model.h5')
```

## Improving the model with Data augmentation and Dropout

```
In [10]: # create ImageDataGenerator that augments image
# specify to rescale pixel values to [0,1]
train_datagen_augmented = ImageDataGenerator(rescale = 1./255,
                                              # the generator can choose any value within these ranges
                                              rotation_range=40,
                                              width_shift_range=0.2,
                                              height_shift_range=0.2,
                                              shear_range=0.2,
                                              zoom_range=0.2,
                                              horizontal_flip=True)

# data augmentation NOT required in validation set
test_datagen = ImageDataGenerator(rescale = 1./255)

# this generator generates data # directory variable specified earlier
train_generator_augmented = train_datagen_augmented.flow_from_directory(train_dir,
                                # photos are resized to 150 x 150
                                target_size = (150,150),
                                # IMPT: images are generated in batches
                                batch_size = 20,
                                # specify binary classification
                                class_mode = 'binary')

# directory variable specified earlier
validation_generator = test_datagen.flow_from_directory(validation_dir,
                                                        # photos are resized to 150 x 150
                                                        target_size = (150,150),
                                                        # photos will come in batches of 20
                                                        batch_size = 20,
                                                        # specify binary classification
                                                        class_mode = 'binary')
```

Found 2000 images belonging to 2 classes.  
Found 1000 images belonging to 2 classes.

```

In [11]: # model_2 with dropout 50%
model_2 = models.Sequential()
            # using 32 filters of size 3x3 NOTE: default stride=1, default
model_2.add(layers.Conv2D( 32, (3,3), activation='relu',
            # images were resized to 150x150 earlier
            # 3 because images are colored
            input_shape=(150,150,3)))
# each convolution layer has its own max pooling layer
model_2.add(layers.MaxPooling2D((2,2)))

            # 64 filters of size 3x3
model_2.add(layers.Conv2D( 64, (3,3), activation='relu',input_shape=(150,150,3)))
model_2.add(layers.MaxPooling2D((2,2)))

            # 128 filters of size 3x3
model_2.add(layers.Conv2D( 128, (3,3), activation='relu',input_shape=(150,150,3)))
model_2.add(layers.MaxPooling2D((2,2)))

            # 128 filters of size 3x3
model_2.add(layers.Conv2D( 128, (3,3), activation='relu',input_shape=(150,150,3)))
model_2.add(layers.MaxPooling2D((2,2)))

# flatten to single array
model_2.add(layers.Flatten())

# in every epoch 50% of neurons are deactivated to reduce interneuron dependency and ov
model_2.add(layers.Dropout(0.5))

# 512 neuron hidden layer
model_2.add(layers.Dense(512, activation='relu'))

# 1 neuron output layer for binary classification
# sigmoid activation for binary classification
model_2.add(layers.Dense(1, activation='sigmoid'))

```

```

In [12]: model_2.compile(loss='binary_crossentropy',
            # RMSprop used over sgd # Learning rate set to 0.0001
            optimizer= optimizers.RMSprop(lr=1e-4),
            metrics=['acc'])

```

In [13]:

```
# model is fit into augmented data
history = model_2.fit_generator(train_generator_augmented,
                                # IMPT: 32 x 100 = 3200 augmented images
                                steps_per_epoch=100,
                                epochs=20,
                                # validation data comes from the validation generator
                                validation_data=validation_generator,
                                # IMPT 32 x 50 = 1600 images in the validation set
                                validation_steps = 50)
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...
to
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...
to
['...']
```

Train for 100 steps, validate for 50 steps

Epoch 1/20

100/100 [=====] - 76s 763ms/step - loss: 0.6933 - acc: 0.5

150 - val\_loss: 0.6851 - val\_acc: 0.5430

Epoch 2/20

100/100 [=====] - 66s 663ms/step - loss: 0.6848 - acc: 0.5

510 - val\_loss: 0.6910 - val\_acc: 0.5040

Epoch 3/20

100/100 [=====] - 80s 803ms/step - loss: 0.6782 - acc: 0.5


720 - val\_loss: 0.6665 - val\_acc: 0.5680

Epoch 4/20

100/100 [=====] - 66s 663ms/step - loss: 0.6800 - acc: 0.5

In [14]:

```
# visualise the training progression
pd.DataFrame(history.history).plot(figsize=(8,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.show()
```



- clearly, with the augmentation and dropot, overfitting is much lesser

```
In [15]: # save model
model_2.save('CNN_model_augmentation_plus_dropout.h5')
```

## Improving mode with Transfer Learning

```
In [16]: # create ImageDataGenerator that augments image
# specify to rescale pixel values to [0,1]
train_datagen_augmented = ImageDataGenerator(rescale = 1./255,
                                              # the generator can choose any value within these ranges
                                              rotation_range=40,
                                              width_shift_range=0.2,
                                              height_shift_range=0.2,
                                              shear_range=0.2,
                                              zoom_range=0.2,
                                              horizontal_flip=True)

# data augmentation NOT required in validation set
test_datagen = ImageDataGenerator(rescale = 1./255)

# this generator generates data          # directory variable specified earlier
train_generator_augmented = train_datagen_augmented.flow_from_directory(train_dir,
                                                                           # photos are resized to 150 x 150
                                                                           target_size = (150,150),
                                                                           # IMPT: images are generated in batches
                                                                           batch_size = 20,
                                                                           # specify binary classification
                                                                           class_mode = 'binary')

# directory variable specified earlier
validation_generator = test_datagen.flow_from_directory(validation_dir,
                                                         # photos are resized to 150 x 150
                                                         target_size = (150,150),
                                                         # photos will come in batches of 20
                                                         batch_size = 20,
                                                         # specify binary classification
                                                         class_mode = 'binary')
```

Found 2000 images belonging to 2 classes.  
Found 1000 images belonging to 2 classes.

```
In [17]: # convotional base of VGG16 is used instead of creating our own
          # these weights are the VGG16 final weights from imagenet competition
conv_base = VGG16(weights = 'imagenet',
                  # specify that we only want the convolution base
                  include_top = False,
                  # colored input was resized to 150 x 150 earlier
                  input_shape = (150,150,3))
```



In [18]: conv\_base.summary()

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 150, 150, 3)]	0
-----		
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
-----		
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
-----		
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
-----		
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
-----		
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
-----		
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
-----		
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168

In [19]: *# before trainging, decide whether to let the convolutoon layer be trainable*  
conv\_base.trainable = **False**

In [20]: *# initialise the model w VGG16 convolution base*  
model\_3 = models.Sequential()  
model\_3.add(conv\_base)  
  
*# flatten to single array*  
model\_3.add(layers.Flatten())  
  
*# 256 neuron hidden layer*  
model\_3.add(layers.Dense(256, activation='relu'))  
  
*# 1 neuron output layer for binary classification*  
*# sigmoid activation for binary classification*  
model\_3.add(layers.Dense(1, activation='sigmoid'))

In [21]: *# by freezing the convolution base, there are 14 mil less params to train*  
model\_3.summary()

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
vgg16 (Model)	(None, 4, 4, 512)	14714688
-----		
flatten_2 (Flatten)	(None, 8192)	0
-----		
dense_4 (Dense)	(None, 256)	2097408
-----		
dense_5 (Dense)	(None, 1)	257
=====		
Total params: 16,812,353		
Trainable params: 2,097,665		
Non-trainable params: 14,714,688		
-----		

```
In [19]: conv_base.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 150, 150, 3)]	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

```
In [22]: model_3.compile(loss='binary_crossentropy',  
                        # RMSprop used over SGD # Learning rate set much lower because we are  
                        optimizer= optimizers.RMSprop(lr=2e-5),  
                        metrics=['acc'])
```

```
In [23]: # w early stopping we can set epochs to a high number and training will stop itself
checkpoint_cb = keras.callbacks.ModelCheckpoint('CNN_VGG16_early_stopping_model.h5', save_best_only=True)
early_stopping_cb = keras.callbacks.EarlyStopping(patience = 10,
                                                    restore_best_weights = True)
```

```
In [37]: # model is fit into augmented data
history = model_3.fit_generator(train_generator_augmented,
                                # IMPT: 32 x 100 = 3200 augmented images
                                steps_per_epoch=100,
                                epochs=2,
                                # validation data comes from the validation generator
                                validation_data=validation_generator,
                                # IMPT 32 x 50 = 1600 images in the validation set
                                validation_steps = 50,
                                # specify to save checkpoint via early stopping
                                callbacks=[checkpoint_cb, early_stopping_cb])
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...
to
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...
to
['...']
```

Train for 100 steps, validate for 50 steps


Epoch 1/2

100/100 [=====] - 954s 10s/step - loss: 0.2941 - acc: 0.8750  
- val\_loss: 0.2318 - val\_acc: 0.8990

Epoch 2/2

100/100 [=====] - 908s 9s/step - loss: 0.2832 - acc: 0.8740 -  
val\_loss: 0.2333 - val\_acc: 0.8990

```
In [38]: # 31th to 32th epoch
pd.DataFrame(history.history).plot(figsize=(8,5))
plt.grid(True)
plt.gca().set_ylim(0,1)
plt.show()
```



```
In [34]: # Load saved model
model_3 = keras.models.load_model('CNN_VGG16_early_stopping_model.h5')
```

- validation has **plateau'ed at 0.90 with very little overfitting**
- **setting convolution base to trainable will improve accuracy significantly**, but will also take more time due to more params